

# Data Science for Business

## Lecture #4

### *Topic Modeling for Movie Reviews Problem*

---

**Prof. Alan L. Montgomery**

The University of Hong Kong & Carnegie Mellon University, Tepper School of Business

email: [alanmontgomery@cmu.edu](mailto:alanmontgomery@cmu.edu)

All Rights Reserved, © 2021 Alan Montgomery

Do not distribute, post, or reproduce without Alan Montgomery's Permission

# Outline

---

Clustering Movies using Text from Movie Reviews

Small Example for Topic Modeling applied to Movie Reviews




# Clustering Movies using the Text from Movie Reviews


---

Need to structure “unstructured data/big data”






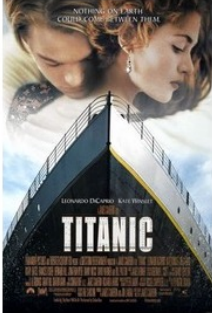
# How do we analyze user reviews?



Search movies, TV, actors, more... 


MOVIES & DVDS ▾ TV ▾ NEWS ▾ TICKETS & SHOWTIMES

TRENDING ON RT June's Most Anticipated Movies Future of DC Movies 100 Best Sci-Fi TV Shows June TV Binge Guide   



**TITANIC**

PG-13, 3 hr. 14 min.  
[Drama, Romance](#)  
Directed By: [James Cameron](#)  
In Theaters: Dec 19, 1997  
On DVD: Sep 10, 2012  
Paramount Pictures




**Titanic**  
2 minutes 17 seconds  
Added: May 9, 2008  
[View All Videos \(1\) >](#)

## TITANIC REVIEWS

All Critics Top Critics My Critics DVD Audience

< Page 1 of 69743 >




**Leroy G**

★★★★½

June 3, 2018

A tragic romantic movie I watched for the first time with my girlfriend. Over twenty years later we still watch it together.




**Tiffany H**

★★★★★

June 3, 2018

titanic is one of the best romance movie I've ever watch. cinematography, story, acting is realistic and classic.



**Brendan N** ★ Super Reviewer

★★★★★

May 28, 2018

James Cameron might have seemed the incorrect choice to direct this film but he takes a step back from action films to deliver one of the better epics. The film is strengthened by Winslet, she alone carries the better half of the film, her arc is probably the strongest. Billy Zane is excellent as the jealous and conniving husband, while the ice berg might steal the show as the core villain, Billy is doing all the right things as the human villain. DiCaprio was at the beginning of his career and without this cementing him in the hearts of every person in the world, there wouldn't be the long career he has been granted. He isn't the best in the lead, he lacks that presence but to be honest, he doesn't have a lot to play with. Did the film deserve all the Oscar glory? Probably not. the filmmaking is flawless but the story and plotting are the weakest parts. It was a great film because Hollywood ignores these sweeping epics, but look at the box office, some gambles are worth taking. 25/05/2018.



# An approach for structuring is to map text to “tags”

*Tags can be user generated like on Flickr or del.icio.us or tagged automatically*

---

James Cameron might have seemed the incorrect choice to direct this film but he takes a step back from action films to deliver one of the better epics. The film is strengthened by Winslet, she alone carries the better half of the film, her arc is probably the strongest. Billy Zane is excellent as the jealous and conniving husband, while the iceberg might steal the show as the core villain, Billy is doing all the right things as the human villain. DiCaprio was at the beginning of his career and without this cementing him in the hearts of every person in the world, there wouldn't be the long career he has been granted. He isn't the best in the lead, he lacks that presence but to be honest, he doesn't have a lot to play with. Did the film deserve all the Oscar glory? Probably not. the filmmaking is flawless but the story and plotting are the weakest parts. It was a great film because Hollywood ignores these sweeping epics, but look at the box office, some gambles are worth taking. 25/05/2018.



- James Cameron
- action
- Winslet
- Billy Zane
- iceberg
- villain
- DiCaprio
- Oscar
- flawless
- epic



# User generated “Tags” can be used to understand movies

Example of Information from MovieLens about Titanic summed across users

	A	B	C				
1	odid	tag	count				
2	10100	romance	68	21	10100	shipwreck	9
3	10100	leonardo dicaprio	46	22	10100	survival	8
4	10100	atmospheric	29	23	10100	oscar (best directing)	6
5	10100	disaster	27	24	10100	action	5
6	10100	love story	27	25	10100	history	5
7	10100	true story	25	26	10100	music	5
8	10100	drama	23	27	10100	overrated	5
9	10100	kate winslet	23	28	10100	time travel	5
10	10100	bittersweet	22	29	10100	love	4
11	10100	oscar (best picture)	22	30	10100	kathy bates	3
12	10100	historical	21	31	10100	70mm	2
13	10100	catastrophe	20	32	10100	big budget	2
14	10100	chick flick	15	33	10100	class differences	2
15	10100	based on a true story	14	34	10100	epic	2
16	10100	sentimental	14	35	10100	girlie movie	2
17	10100	james cameron	12	36	10100	natural disaster	2
18	10100	nudity (topless - notable)	11				
19	10100	nudity (topless)	10				
20	10100	oscar (best cinematography)	9				

These are our **Words**. Each document only contains a subset of all the words in the vocabulary (e.g., Titanic uses 47 out of 962 terms). But these words tend to group together (e.g., “romance” and “love” tend to go together).



This yields a *huge* matrix of 962 terms by 1,153 movies

```
> mat_dtm[1:15,toptags]
```

Docs	Terms									
	sci-fi	action	funny	visually appealing	superhero	animation	comedy	based on a book	predictable	twist ending
Titanic	0	5	0	0	0	0	0	0	1	0
The Dark Knight	0	57	0	0	97	0	0	0	0	0
Star Wars Ep. I: The P	42	4	0	1	0	0	0	0	0	0
Pirates of the Caribbe (2006)	0	8	7	0	0	0	14	0	0	0
Transformers: Revenge	3	12	1	0	0	0	2	0	0	0
Jurassic Park	58	32	5	0	0	0	0	21	1	0
Finding Nemo	0	0	22	0	0	51	10	0	4	0
Spider-Man 3	0	5	0	0	18	0	0	0	0	0
The Lion King	0	0	0	0	0	44	0	0	0	0
Shrek the Third	0	0	0	0	0	2	2	0	0	0
Transformers	10	9	0	0	0	0	1	0	0	0
Iron Man	29	23	21	9	74	0	2	0	1	0
Indiana Jones and the	0	1	0	0	0	0	1	0	0	0
Pirates of the Caribbe (2007)	0	2	2	0	0	0	13	0	0	0
Harry Potter and the H	0	0	6	0	0	0	11	13	0	0

Our objective is to reduce the dimensionality of the matrix and find interesting patterns. In other words group the words/rows and also group the movies/columns. Could use SVD which is often used for recommender systems.





## Data: Movie Tags

```
> as.matrix(mterms[1:15,topterms])
```

	animation	based on a book	comedy	funny	nudity (topless)	predictable	remake
Titanic	0	0	0	0	10	1	0
The Dark Knight	0	0	0	0	0	0	0
Star Wars Ep. I: The	0	0	0	0	0	0	0
Pirates of the Carib	0	0	14	7	0	0	0
Transformers: Reveng	0	0	2	1	0	0	0
Jurassic Park	0	21	0	5	0	1	0
Finding Nemo	51	0	10	22	0	4	0
Spider-Man 3	0	0	0	0	0	0	0
The Lion King	44	0	0	0	0	0	0
Shrek the Third	2	0	2	0	0	0	0
Transformers	0	0	1	0	0	0	0
Iron Man	0	0	2	21	0	1	0
Indiana Jones and th	0	0	1	0	0	0	0
Pirates of the Carib	0	0	13	2	0	0	0
Harry Potter and the	0	13	11	6	0	0	0





# Relative Frequency of Words

```
> round(mtxterms[1:15,topterms],2)
```

	animation	based on a book	comedy	funny	nudity (topless)	predictable	remake
Titanic	0.00	0.00	0.00	0.00	0.02	0.00	0
The Dark Knight	0.00	0.00	0.00	0.00	0.00	0.00	0
Star Wars Ep. I: The	0.00	0.00	0.00	0.00	0.00	0.00	0
Pirates of the Carib	0.00	0.00	0.06	0.03	0.00	0.00	0
Transformers: Reveng	0.00	0.00	0.01	0.01	0.00	0.00	0
Jurassic Park	0.00	0.03	0.00	0.01	0.00	0.00	0
Finding Nemo	0.15	0.00	0.03	0.06	0.00	0.01	0
Spider-Man 3	0.00	0.00	0.00	0.00	0.00	0.00	0
The Lion King	0.15	0.00	0.00	0.00	0.00	0.00	0
Shrek the Third	0.04	0.00	0.04	0.00	0.00	0.00	0
Transformers	0.00	0.00	0.00	0.00	0.00	0.00	0
Iron Man	0.00	0.00	0.00	0.04	0.00	0.00	0
Indiana Jones and th	0.00	0.00	0.01	0.00	0.00	0.00	0
Pirates of the Carib	0.00	0.00	0.06	0.01	0.00	0.00	0
Harry Potter and the	0.00	0.06	0.05	0.03	0.00	0.00	0



# What problems can we address with our review data?

---

## Movie recommendation system

- What is another movie that user hasn't viewed that they might like?
- Look for similar movies

## Planning first-run release schedule

- What is the best week to release a movie?
- Look for weeks in which other (announced) releases are dis-similar, in other words find less competition

## Dimensionality reduction

- Create a new embedding space that compresses our high-dimensional tag vector into a low dimension numeric space that preserves as much information as possible



## Clustering Movies with Topic Models

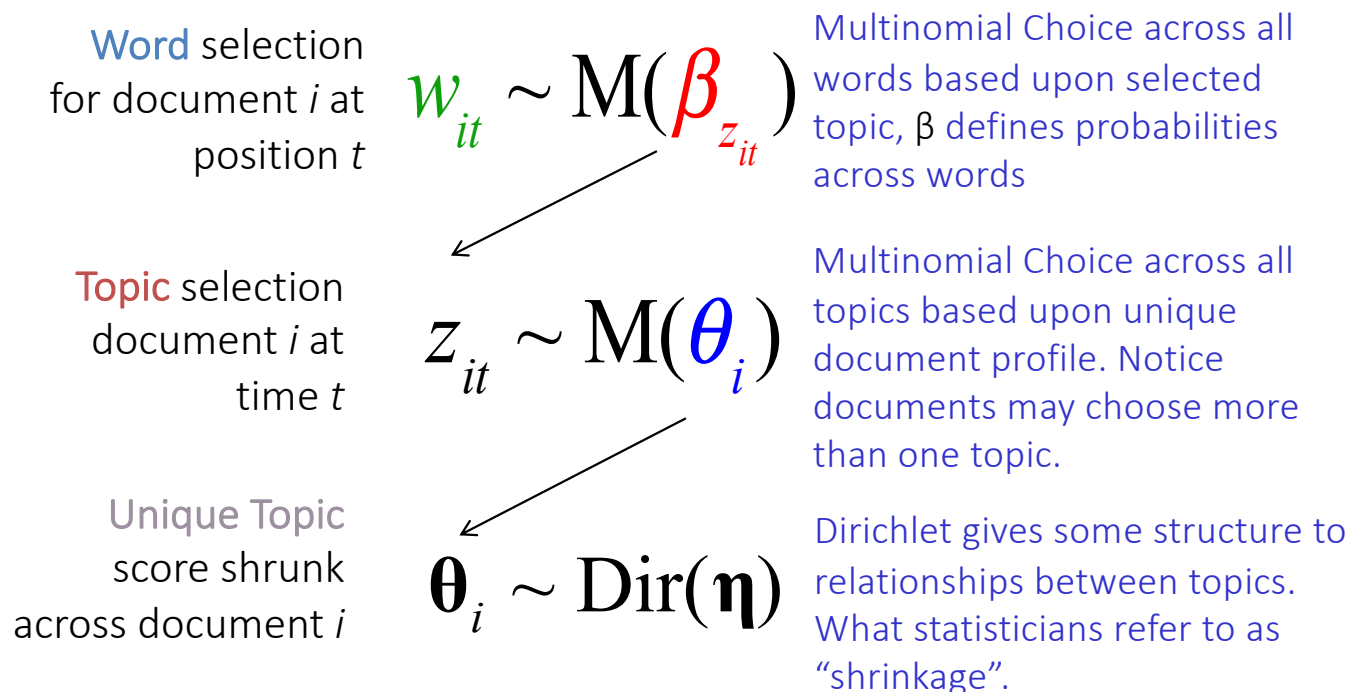
---

Topic models or Latent Dirichlet Allocation models are a type of probabilistic clustering algorithm well suited for sparse text data



# Latent Dirichlet Allocation detects “topics” which provide latent dimensions to describe observations

---



# Understanding our topic model with a simple example

Suppose we observe the following three movies

*What do our topics mean?*

Our LDA model learns the probability of a tag given a topic:

*We can use unique topic profile to compare movies*

Also our LDA model learns the distribution of topics associated with each movie:

Movie	“Adventure”	“Funny”
Jurassic Park	50	0
Iron Man	10	10
Finding Nemo	5	25

This is our raw data that has the counts of the tags

Tag	Topic1	Topic2
Adventure	100%	10%
Funny	0	90%

Each column corresponds with  $\beta$  or  $\text{Prob}(\text{Word} | \text{Topic})$

Movie	Topic1	Topic2
Jurassic Park	100%	0
Iron Man	50%	50%
Finding Nemo	0	100%

Each row corresponds with  $\theta$  or  $\text{Prob}(\text{Topic})$  for a movie



# Predicting the word count by multiplying tag distribution for each topic by the topic distribution

Our LDA model learns the probability of a tag given a topic:

Tag	Topic1	Topic2
Adventure	100%	10%
Funny	0	90%

Each column corresponds with  $\beta$  or  $\text{Prob}(\text{Word} | \text{Topic})$

The unique distribution of topics to Iron Man:

Movie	Topic1	Topic2
Iron Man	50%	50%

Each row corresponds with  $\theta$  or  $\text{Prob}(\text{Topic})$  for a movie

To construct predictions of the words counts we can multiply the conditional probability of words given the tags by the movies topic profile by word count:

Movie	"Adventure"	"Funny"
Topic 1	100% x 50%	0% x 50%
Topic 2	10% x 50%	90% x 50%
Total (weight by word count)	$= (0.50 + 0.05) \times 20$ $= 0.55 \times 20 = 11$	$= (0 + 0.45) \times 20$ $= 0.45 \times 20 = 9$

Assumes we know there are 20 tags in total:  
we predict 11 Adventure and 9 Funny

This process illustrates how we move from the parameters to make inferences about the data. When training this model we make use of Bayes rule to invert the process and infer the parameters given the data.



## Topic Modeling applied to the Movie Data with a Small Scale Dataset

---

See “movie\_example.zip”





# Setup and Read in the Data

First we need to load in the libraries that are needed. Topic Model often requires large matrices to be represented, but instead of representing this as a “dense” matrix we only store the values that are set using a “sparse” matrix

```
23 ##### setup environment #####
24
25 # setup libraries
26 if (!require(lattice)) {install.packages("lattice"); library(lattice)}
27 if (!require(NLP)) {install.packages("NLP"); library(NLP)}
28 if (!require(topicmodels)) {install.packages("topicmodels"); library(topicmodels)}
29 if (!require(tm)) {install.packages("tm"); library(tm)}
30 if (!require(slam)) {install.packages("slam"); library(slam)}
31
32
33
34
35 ##### input the data #####
36
37 ## read in the data
38
39 # in RStudio select Menu Bar --> Session --> Set Working Directory --> To Source File Directory
40 # or automatically set working directory to be that of the script
41 setwd(dirname(rstudioapi::getActiveDocumentContext()$path)) # only works in Rstudio scripts
42 # alternatively set the working directory manually
43 #setwd("~/Documents/class/marketing analytics/cases/movies") # !! edit and uncomment this line, if needed !!
44
45 # read in movie datasets
46 movies=read.delim("opus_movies.txt",header=T) # the Opus movie data
47 tags=read.delim("opus_movielsens_tags.txt",header=T) # just the tags from movielens
48
49
50 ## make modifications to the dataset
51
52 # change data formats
53 tags$odid=as.factor(tags$odid)
```



# Transform the Data

I'll only save movies produced by Paramount and that use the three terms “action”, “comic book” and “animation”.

Our mterms matrix is a sparse matrix. You can reference the rows and columns using the selection as usual, but the result is another sparse matrix. If you want to work with the matrix in the usual way cast it as follows:

```
> as.matrix(mterms[1:2,1:2])
      action comic book
220100      23      31
5580100      0       0
> umovies$display_name[umovies$odid %in% c(220100,5580100)]
[1] Iron Man Beowulf
```

```
56 ## transform the terms into a structure that can be used for topic modeling
57
58 # use this definition of mterms for movielens tags
59 # put data in sparse matrix form using simple_triplet_matrix as needed by LDA
60 mterms=simple_triplet_matrix(i=as.integer(tags$odid),j=as.integer(tags$tag),v=tags$count,
61                             dimnames=list(levels(tags$odid),levels(tags$tag)))
62 # let's create a list of a smaller list of movies produced by paramount
63 movielist=movies$odid[movies$production_company1=="Paramount Pictures"]
64 # let's create a short list of terms to save
65 shorttermList=c("action","comic book","animation")
66 # keep only the subset of a few terms
67 mterms=mterms[rownames(mterms) %in% movielist,shorttermList]
68 # also delete any movies that do not have any terms
69 mterms=mterms[apply(mterms,1,sum)>0,]
70 # let's update the list of movies and their names since some might have just been deleted
71 movielist=rownames(mterms)
72 movienames=as.character(movies$display_name[movies$odid %in% rownames(mterms)])
73 # let's print out our matrix
74 as.matrix(mterms)
75 # let's lookup the movie names
76 movies[movies$odid %in% rownames(mterms),c("odid","display_name")]
77 # compute totals for mterms
78 lterms=apply(mterms,1,sum) # compute the sum of each of the rows (# of terms per movie)
79 lwterms=apply(mterms,2,sum) # compute the sum of each of the columns (# of times word used)
80
81 # prepare a subset of movies with just the movies in our list
82 umovies=movies[movies$odid %in% as.integer(movielist),] # create a subset of the movies that have terms
```



# Perform a k-Means for Comparison

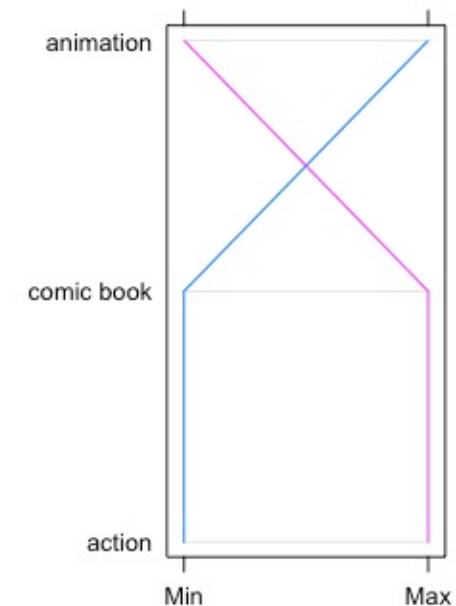
```

87 ##### for reference compute kmeans cluster #####
88
89 # estimate kmeans with two topics
90 (grpKmeans=kmeans(mterms,centers=2))
91
92 # summarize the centroids
93 grpKcenter=t(grpKmeans$centers)
94 parallelplot(t(grpKcenter))
95
96 # print a table with the movies assigned to each cluster
97 for (i in 1:2) {
98   print(paste("* * * Movies in Cluster #",i," * * *"))
99   print(movienames[grpKmeans$cluster==i])
100 }

```

> results=cbind(as.matrix(mterms),grpKmeans\$cluster)  
 > rownames(results)=umovies\$display\_name[rownames(results)%in%umovies\$odid]  
 > results

	action	comic book	animation
Iron Man	23	31	0 2
Beowulf	0	0	8 1
Transformers: Dark of the Moon	7	0	0 1
Puss in Boots	0	0	6 1
Iron Man 2	16	17	0 2
Captain America: The First Avenger	12	14	0 2
Rango	0	0	14 1
The Adventures of Tintin	8	12	21 1
True Grit	2	0	0 1
G.I. Joe: Retaliation	6	0	0 1
Hansel & Gretel: Witch Hunters	3	0	0 1
Jack Reacher	9	0	0 1



# Train our LDA Topic Model using LDA

---

```
104 ▾ ##### estimate an LDA topic model using keywords #####
105
106 ## our first step is to estimate the topic model using LDA
107
108 # setup the parameters for LDA control vector
109 burnin=1000      # number of initial iterations to discard for Gibbs sampler (for slow processors use 500)
110 iter=5000       # number of iterations to use for estimation (for slow processors use 1000)
111 thin=50         # only save every 50th iteration to save on storage
112 seed=list(203,5,63,101,765) # random number generator seeds
113 nstart=5        # number of repeated random starts
114 best=TRUE       # only return the model with maximum posterior likelihood
115
116 # estimate a series of LDA models (each run can take a few minutes depending upon your processor)
117 ClusterOUT = LDA(mterms,2,method="Gibbs",control=list(nstart=nstart,seed=seed,best=best,burnin=burnin,iter=iter,thin=thin))
```

Caution: With a large dataset and a large number of topics this may take hours (or days). Test with small values first.



# Output from LDA gives us Prob(Topic) for each movie and Prob(Word|Topic) for each topic

---

```
120 ## now that we have saved the LDA results to our ClusterOUT object we want to
121 ## extract the topic information and look at them
122
123 # probability of topic assignments (each movie has its own unique profile)
124 # rows are movies and columns are topics
125 ClustAssign = ClusterOUT@gamma # this is a matrix with the row as the movie and column as the topic
126 rownames(ClustAssign)=movienames # set the movie titles as the row names
127 dim(ClustAssign) # check the dimension of the cluster (movies X topics)
128 head(ClustAssign,n=10) # show the actual topic probabilities associated with the first 10 movies
129
130 # matrix with probabilities of each term per topic
131 ClustTopics = exp(ClusterOUT@beta) # notice that we use "@" to access elements in the object and not "$" since this is an S4 object
132 colnames(ClustTopics)=colnames(mterms) # the columns are the terms
133 dim(ClustTopics) # check dimensions of the topics
134 print(ClustTopics) # print out clusters (topics in rows and terms in columns)
```





# Movies are Mixtures of Topics

```

137 ## let's work on understanding the cluster based upon the movies
138
139 # visualize the distribution of topics across the movies
140 boxplot(ClustAssign,xlab="Topic",ylab="Probability of Topic across Movies")
141
142 # print a table with the movies assigned to each cluster
143 ClustBest = apply(ClustAssign,1,which.max) # determine the best guess of a cluster, a vector with best guess
144 for (i in 1:2) {
145     print(paste("* * * Movies in Cluster #",i," * * *"))
146     print(movienames[ClustBest==i])
147 }

```

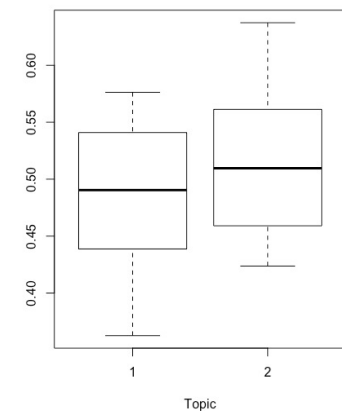
> ClustAssign

	[,1]	[,2]
Iron Man	0.4615385	0.5384615
Beowulf	0.4310345	0.5689655
Transformers: Dark of the Moon	0.5614035	0.4385965
Puss in Boots	0.4464286	0.5535714
Iron Man 2	0.4939759	0.5060241
Captain America: The First Avenger	0.4868421	0.5131579
Rango	0.3906250	0.6093750
The Adventures of Tintin	0.3626374	0.6373626
True Grit	0.5192308	0.4807692
G.I. Joe: Retaliation	0.5535714	0.4464286
Hansel & Gretel: Witch Hunters	0.5283019	0.4716981
Jack Reacher	0.5762712	0.4237288

```

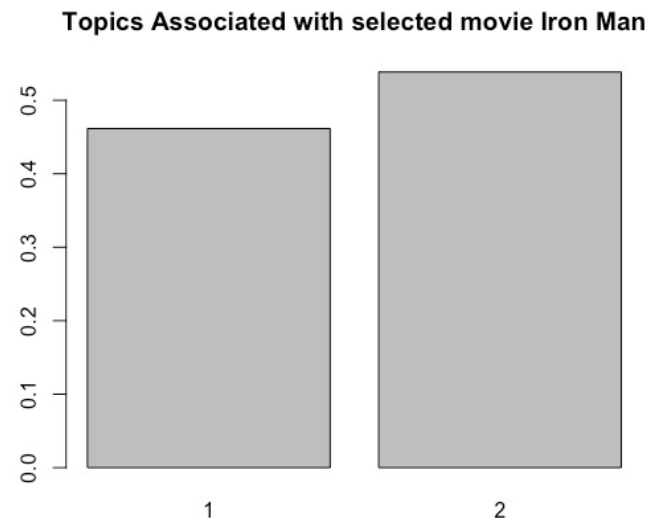
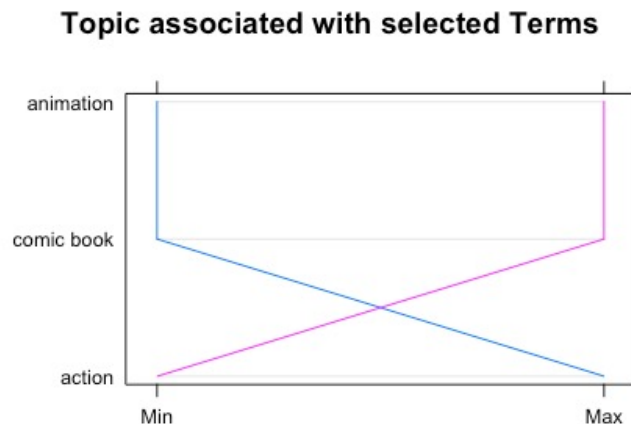
[1] "* * * Movies in Cluster # 1 * * *"
[1] "Transformers: Dark of the Moon" "True Grit"
[3] "G.I. Joe: Retaliation"          "Hansel & Gretel: Witch Hunters"
[5] "Jack Reacher"
[1] "* * * Movies in Cluster # 2 * * *"
[1] "Iron Man"
[2] "Beowulf"
[3] "Puss in Boots"
[4] "Iron Man 2"
[5] "Captain America: The First Avenger"
[6] "Rango"
[7] "The Adventures of Tintin"

```



# To understand a Topic look at its Associations

```
150 ## another way to understand the topics is through their associations with the keywords
151
152 # show the terms and associated topics
153 parallelplot(ClustTopics,main="Topic associated with selected Terms")
154
155 # show the topics associated with a selected movie
156 imovie=1
157 barplot(ClustAssign[imovie,],names.arg=1:ncol(ClustAssign),main=paste("Topics Associated with selected movie",umovies$display_name[imovie]))
```





# We can predict the Words using our Model

```
160 ## last we can use our model to compute a best guess
161
162 # determine the best guess for each movie/term combination
163 ClustGuess=(ClustAssign%*%ClustTopics)*lmterms
164
165 # we can compare the predictions for a selected movie
166 imovie=1
167 mcompare=cbind(ClustGuess[imovie,],as.vector(mterms[imovie,]))
168 print(mcompare)
169
170 # or we can print the predictions for all movies
171 as.matrix(cbind(ClustGuess,mterms))
172
173 # compare kmeans solutions with the topic model
174 # remember that kmeans assignments are deterministic, while topic models are probabilistic
175 # so this cross tab only considers the matches between the most likely
176 xtabs(~grpKmeans$cluster+ClustBest)
```

> # or we can print the predictions for all movies

> as.matrix(cbind(ClustGuess,mterms))

	action	comic book	animation	action	comic book	animation
Iron Man	24.888900	17.503332	11.607767	23	31	0
Beowulf	3.443976	2.739460	1.816563	0	0	8
Transformers: Dark of the Moon	3.923207	1.849646	1.227145	7	0	0
Puss in Boots	2.675058	1.999193	1.325748	0	0	6
Iron Man 2	16.276970	10.054418	6.668612	16	17	0
Captain America: The First Avenger	12.639381	8.032915	5.327703	12	14	0
Rango	5.462995	5.133389	3.403614	0	0	14
The Adventures of Tintin	14.854869	15.721781	10.423349	8	12	21
True Grit	1.036835	0.579062	0.384103	2	0	0
G.I. Joe: Retaliation	3.315904	1.613598	1.070498	6	0	0
Hansel & Gretel: Witch Hunters	1.582380	0.852270	0.565349	3	0	0
Jack Reacher	5.177514	2.297856	1.524629	9	0	0



# Conclusion

---



# Summary

---

Instead of clustering movies based upon attributes (like genre, movie stars, budget, ratings, ...) we will cluster them using the text of movie reviews

We take the text of the movie reviews and pre-process them to yield “tags” that describe each movie. In our example, we have almost 1,000 terms that describe over 1,000 movies.

Topic modeling wishes to find “topics” or clusters that explain the kinds of words that we will see.

- The topics are common across movies, and you may find some topics are similar to genres like romance or science fiction, but a difference is that every word has some chance of showing up in a topic. Even “action” has a small chance of occurring in a romance.
- Every movie has a unique profile of topics, and this is what gives us our dimensionality reduction.

Unlike k-Means analysis where each observation is in one cluster, in topic modeling each observation has a chance of being assigned to every topic.

